

Recitation 9: Tshlab (& VM)

Instructor: TAs

28 March 2019

Outline

- Labs
- Signals
- IO

tshlab and malloclab

- **tshlab due Thursday, March 28. Start Early!!**

- **malloclab is released tomorrow**
 - Start early
 - Do the checkpoint first, don't immediately go for the final
 - Expect a recitation next week
 - Working for several hours will improve the value significantly

Signals

- **Parent process sends SIGINT to a child process.
What is the behavior of the child?**
- **What is the default?**
- **What else could the child do?**

More Signals

- **Parent process sends SIGKILL to a child process.
What is the behavior of the child?**
- **What is the default?**
- **What else could the child do?**

Sending Signals

- Parent sends SIGKILL to a child process.

...

```
pid_t pid = ...; // child pid
kill(pid, SIGKILL);
// At this point, what could have
// happened to the child process?
```

Blocking Signals

- The shell is currently running its handler for SIGCHLD.
- What signals can it receive?
- What signals can it not receive (i.e., blocked)?

Errno

- Included from `<errno.h>`
- Global integer variable – usually 0
- When a system call fails (usually indicated by returning -1), it also will set `errno` to a value describing what went wrong

- Example: let's assume there is no “foo.txt” in our path

```
int fd = open("foo.txt", O_RDONLY);  
if (fd < 0) printf("%d\n", errno);
```

- What would the code above print?
- The code above will print 2 – in the man pages, we can see that 2 is `ENOENT` “No such file or directory”
- In shell lab, your signal handlers must preserve `errno`

Files

Needed for tshlab

- `int open(const char *pathname, int flags);`
- `int close(int fd);`
- `int dup2(int oldfd, int newfd);`

Needed for life

- `ssize_t read(int fd, void *buf, size_t count);`
- `ssize_t write(int fd, const void *buf, size_t count);`
- `off_t lseek(int fd, off_t offset, int whence);`

IO functions

Needed for tshlab

- `int open(const char *pathname, int flags);`
- `int close(int fd);`
- `int dup2(int oldfd, int newfd);`


Needed for life

- `ssize_t read(int fd, void *buf, size_t count);`
- `ssize_t write(int fd, const void *buf, size_t count);`
- `off_t lseek(int fd, off_t offset, int whence);`

More on dup2

```
int dup2(int oldfd, int newfd)
```

- Makes newfd be the copy of oldfd, closing newfd first if necessary.
- oldfd not valid then newfd not closed and returns -1 errno set
- oldfd == newfd then does nothing and returns newfd




File descriptors

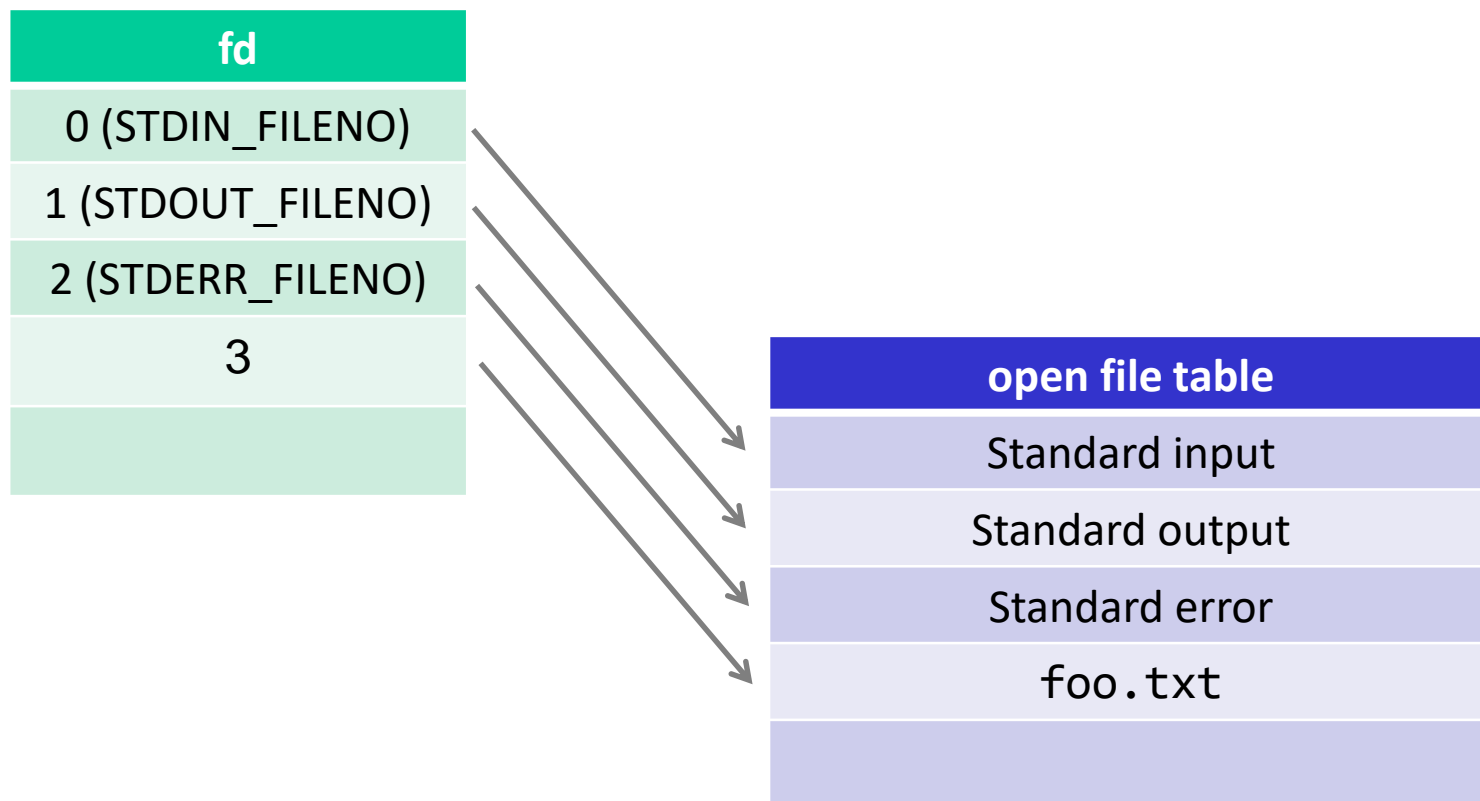
fd
0 (STDIN_FILENO)
1 (STDOUT_FILENO)
2 (STDERR_FILENO)

**stdin, stdout, stderr are
opened automatically
and closed by normal
termination or exit()**

open file table
Standard input
Standard output
Standard error



open("foo.txt")



open("foo.txt")

fd
0 (STDIN_FILENO)
1 (STDOUT_FILENO)
2 (STDERR_FILENO)
3
4

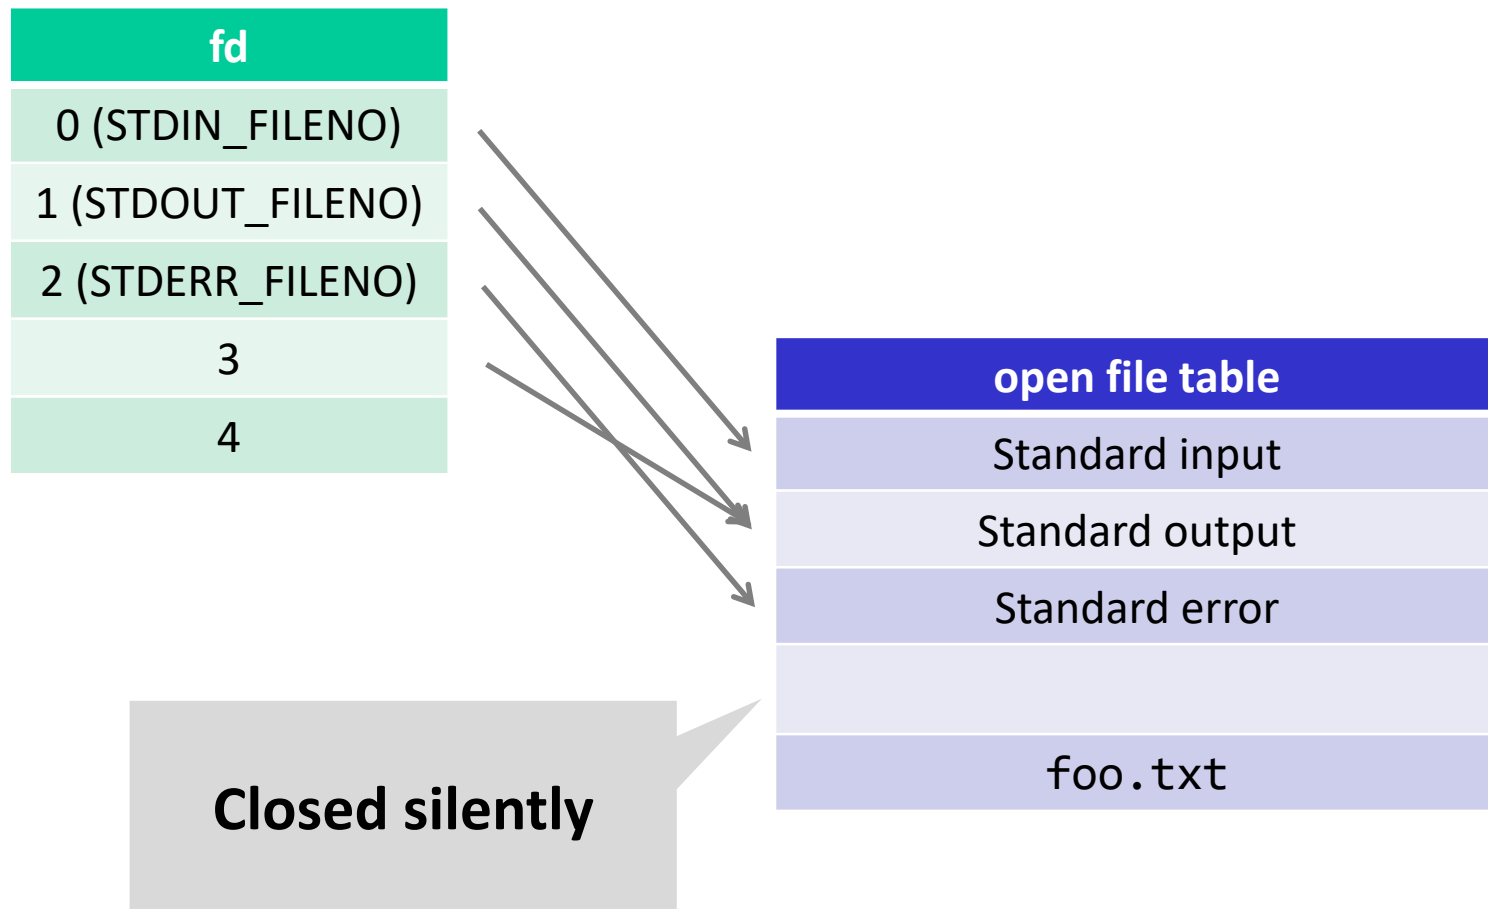
inode table
foo.txt

open file table
Standard input
Standard output
Standard error
foo.txt
foo.txt

**Each call to open()
creates a new open file
description**



dup2(STDOUT_FILENO, 3)



More on open

- `int open(const char *pathname,
 int flags, mode_t mode);`
- For *flags*, you can pass a **bitwise-OR** of one or more flags
- Three kinds of flags (we only discuss the important ones)
 - Access modes (one of them must be included):
 - `O_RDONLY`, `O_WRONLY`, `O_RDWR`
 - File creation flags:
 - `O_CREAT`, `O_TRUNC`, etc.
 - File status flags
- What kind of file does each of the following code opens?
 - `open("foo.txt", O_WRONLY|O_CREAT)`
 - `open("foo.txt", O_RDWR|O_TRUNC)`

More on open

- `int open(const char *pathname,
 int flags, mode_t mode);`
- For *mode*, you can pass a bitwise-OR of one or more constants
- Specifies, when creating a file, what permission the file will be created with
- Only useful when *flags* contain `O_CREAT` (or `O_TMPFILE`)

Linux permissions

- **Every file and directory has permission information**
- **You've seen it before**
 - `ls -l` prints the permissions for each file/directory like:
-rw-r--r-- ... drwxr-xr-x ...
 - `chmod` changes the permissions for files/directories
 - `$ chmod -R 777 /`
 - `$ chmod -R 755 /` (DON'T DO THIS)
 - `$ chmod -X bomb`
- **There are read (R), write (W) and executable (X) permissions for user (USR), group (GRP) and other (OTH)**

Specify permissions in open()

	Read (R)	Write (W)	Executable (X)	All (RWX)
User (USR)	S_IRUSR	S_IWUSR	S_IXUSR	S_IRWXU
Group (GRP)	S_IRGRP	S_IWGRP	S_IXGRP	S_IRWXG
Other (OTH)	S_IROTH	S_IWOTH	S_IXOTH	S_IRWXO

- These constants can be bitwise-OR'd and passed to the third argument of open()
- What does `S_IRWXG | S_IXUSR | S_IXOTH` mean?
- How to create a file which everyone can read from but only the user can write to it or execute it?

Map, Unmap, Launch

- What do map and unmap do?
- What does launch do?
 - How should we tell our eval function of the new execution?
- What are the following codes equivalent to?
 - tsh> map .sh /bin/echo -n
 - tsh> launch 15213.sh
- Functions in tsh_exec that are helpful
 - get_entry: looks up if an extension matched with any saved entry
 - set_exec_entry: create entry for an extension
 - destroy_entry: remove an entry with the extension
 - show_maps: list all the current mappings (entries)
 - maperror : print out errors for invalid input, mapping failures and etc.

Remember to check for errors!!

IO and Fork()

- File descriptor management can be tricky.
- How many file descriptors are open in the parent process at the indicated point?
- How many does each child have open at the call to `execve`?

```
int main(int argc, char** argv)
{
    int i;
    for (i = 0; i < 4; i++)
    {
        int fd = open("foo", O_RDONLY);
        pid_t pid = fork();
        if (pid == 0)
        {
            int ofd = open("bar", O_RDONLY);
            execve(...);
        }
    }
    // How many file descriptors are open in the parent?
```

Redirecting IO

- File descriptors can be directed to identify different open files.

```
int main(int argc, char** argv) {
    int i;
    for (i = 0; i < 4; i++)
    {
        int fd = open("foo", O_RDONLY);
        pid_t pid = fork();
        if (pid == 0)
        {
            int ofd = open("bar", O_WRONLY);
            dup2(fd, STDIN_FILENO);
            dup2(ofd, STDOUT_FILENO);
            execve(...);
        }
    }
    // How many file descriptors are open in the parent?
}
```

Redirecting IO

- At the two points (A and B) in main, how many file descriptors are open?

```
int main(int argc, char** argv)
{
    int i, fd;
    fd = open("foo", O_WRONLY) ;
    dup2(fd, STDOUT_FILENO) ;
    // Point A
    close(fd) ;
    // Point B
    . . .
```

If you get stuck on tshlab

- Read the writeup!
- Do manual unit testing before `runtrace` and `sdriver`!
- Post private questions on piazza!

- Read the man pages on the syscalls.
 - Especially the error conditions
 - What errors should terminate the shell?
 - What errors should be reported?

Memory Access

- The processor tries to write to a memory address.
- List different steps that are required to complete this operation.

Memory Access

- The processor tries to write to a memory address.
- List different steps that are required to complete this operation. (non exhaustive list)
- Virtual to physical address conversion (TLB lookup)
- TLB miss
- Page fault, page loaded from disk
- TLB updated, check permissions
- L1 Cache miss (and L2 ... and)
- Request sent to memory
- Memory sends data to processor
- Cache updated

Address Translation with TLB

- Translate 0x15213, given the contents of the TLB and the first 32 entries of the page table below.
- 1MB Virtual Memory
256KB Physical Memory
4KB page size

2-way
set
associative

Index	Tag	PPN	Valid
0	05	13	1
	3F	15	1
1	10	0F	1
	0F	1E	0
2	1F	01	1
	11	1F	0
3	03	2B	1
	1D	23	0

VPN	PPN	Valid	VPN	PPN	Valid
00	17	1	10	26	0
01	28	1	11	17	0
02	14	1	12	0E	1
03	0B	0	13	10	1
04	26	0	14	13	1
05	13	0	15	18	1
06	0F	1	16	31	1
07	10	1	17	12	0
08	1C	0	18	23	1
09	25	1	19	04	0
0A	31	0	1A	0C	1
0B	16	1	1B	2B	0
0C	01	0	1C	1E	0
0D	15	0	1D	3E	1
0E	0C	0	1E	27	1
0F	2B	1	1F	15	1

man wait

Taken from <http://man7.org/linux/man-pages/man2/wait.2.html>

WAIT(2)

Linux Programmer's Manual

WAIT(2)

NAME

wait, waitpid, waitid - wait for process to change state

SYNOPSIS

```
#include <sys/types.h>
```

```
#include <sys/wait.h>
```

```
pid_t wait(int *wstatus);
```

```
pid_t waitpid(pid_t pid, int *wstatus, int options);
```

```
int waitid(idtype_t idtype, id_t id, siginfo_t *infop, int options);
```

```
/* This is the glibc and POSIX interface; see  
   NOTES for information on the raw system call. */
```

man pages (probably) cover all you need

- **What arguments does the function take?**
 - read SYNOPSIS
- **What does the function do?**
 - read DESCRIPTION
- **What does the function return?**
 - read RETURN VALUE
- **What errors can the function fail with?**
 - read ERRORS
- **Is there anything I should watch out for?**
 - read NOTES
- **Different categories for man page entries with the same name**
- **Looking up man pages online is not an academic integrity violation**

Function arguments

- Should I do `dup2(old, new)` or `dup2(new, old)`?
- Read the man page:

\$ man dup2

SYNOPSIS

```
#include <unistd.h>

int dup(int oldfd);
int dup2(int oldfd, int newfd);
```

Function behavior

- How should I write my format string when I need to print a long double in octals with precision 5 and zero-padded?
- Read the man page:

\$ man printf

DESCRIPTION

Flag characters

The character % is followed by zero or more of the following flags:

#	The value should be converted...
0	The value should be zero padded...
-	The converted value is to be left adjusted...
' '	(a space) A blank should be left before...
+	A sign (+ or -) should always ...

Function return

- What does `waitpid()` return with and without `WNOHANG`?
- Read the man page:

\$ man waitpid

RETURN VALUE

`waitpid()`: on success, returns the process ID of the child whose state has changed; if `WNOHANG` was specified and one or more child(ren) specified by *pid* exist, but have not yet changed state, then 0 is returned. On error, -1 is returned.

Each of these calls sets `errno` to an appropriate value in the case of an error.

Potential errors

- How should I check `waitpid` for errors?
- Read the man page:

\$ man waitpid

ERRORS

ECHILD (for `waitpid()` or `waitid()`) The process specified by *pid* (`waitpid()`) or *idtype* and *id* (`waitid()`) does not exist or is not a child of the calling process. (This can happen for one's own child if the action for **SIGCHLD** is set to **SIG_IGN**. See also the Linux Notes section about threads.)

EINTR **WNOHANG** was not set and an unblocked signal or a **SIGCHLD** was caught; see `signal(7)`.

EINVAL The *options* argument was invalid.

Get advice from the developers

- I sprintf from a string into itself, is this okay?
- Read the man page:

\$ man sprintf

NOTES

Some programs imprudently rely on code such as the following

```
to sprintf(buf, "%s", some_further_text, buf);
```

to append text to buf. However, the standard has explicitly noted that such calls (such as the above) will not produce the expected results. For libc implementations of the functions sprintf() and vsnprintf() when the output was truncated. If libc 2.0.0, they would return 1,